



Concurrent and Distributed Training for Deep Learning Methods

Assignment 1

Introduction to Parallel Deep Neural Networks

Due Date: 23/12/24

T.A in charge: Guy Sudai, guy.sudai@campus.technion.ac.il




Part 1

Brief Background

Please refer to lecture 3, Tutorial 3 and the mandatory reading part of HW0.

Conventions

We are going to implement a simple neuronal network. For this, we will use in our implantation the following conventions:



Layers

The input layer is the 0_{th} layer, and the output later is the L_{th} layer.

The number of layers is $N_L = L + 1$.

The size-layer vector is a vector of length $L + 1$ where element i represents the number of neurons in layer i . In the example above, the size-layer vector is $[2,3,1]$.

We will note $size(l_i)$ as the number of neurons in layer l_i .

Weights

Weights in this neural network implementation are a list of numpy matrices.

Given a neural network with $N_L = L + 1$ layers, and a size-layer vector $[size(l_0), \dots, size(l_L)]$:

The weight list is of list of length L , denoted as $[w_{01}, w_{12}, \dots, w_{(L-1)L}]$,

where w_{ij} is a matrix of shape $(size(l_i), size(l_j))$, which corresponds to the weight matrix between layers i, j in the network.



In the example above, the weight list is: $[w_{01}, w_{12}]$, where $w_{01} = \begin{bmatrix} a & c & e \\ b & d & f \end{bmatrix}$ and $w_{12} = [p \quad q \quad r]$.

Biases

Biases in this neural network implementation is a list of one-dimensional vectors.

Given a neural network with $N_L = L + 1$ layers, and a size-layer vector $[size(l_0), \dots, size(l_L)]$:

Bias list is of length L , denoted as $[b_{01}, b_{12}, \dots, b_{(l-1)l}]$,

where b_{ij} is a one-dimensional vector of size $size(l_j)$, where entry k represents the bias of neuron number k in the j_{th} layer.

In the example above, the bias list is: $[b_{01}, b_{12}]$, where $b_{01} = [0, 1, 2]$ and $w_{12} = [0]$.

Z

For input vector x to layer l_{th} , the output z is defined as follows:

$$z = w_{(l-1)l}^T \cdot x + b_{(l-1)l}$$

Activations

Activations of the l_{th} layer is the operation of activation function on the z output of the same layer. The result from the above calculation is used as the input for the $(l + 1)_{th}$ layer.

Implementation:

You will implement server basic components of the described neural network.

In utils.py file, implement the following:

`def sigmoid(x)`: Calculates the standard sigmoid function. This function outputs $f(x)$.

- Sigmoid is a standard activation function, where $f(x) = \frac{1}{1+e^{-x}}$.

`def sigmoid_prime(x)`: Calculates the derivative function of sigmoid with input x .



`def random_weights(sizes)`: Calculates and returns a list of random xavier initialized numpy arrays of shapes $(size[i], size[i + 1])$ for $0 \leq i < N_L$.

- Look at the end of `utils.py` for xavier initialization implementation.

`def zeros_weights(sizes)`: Calculates and returns a list of zeros numpy arrays of shapes $(size[l_i], size[l_{i+1}])$ for $0 \leq i < N_L$.

`def zeros_biases(list)`: Calculates and returns a list of zeros numpy arrays of size $size(l_i)$ for $0 < i \leq N_L$.

`def create_batches (data, labels, batch_size)`: Creates batches of training data.

Returns a list of batches from the training data, where each batch is of `batch_size` size.

If the length of dataset is not dividable by the `batch_size`, then the last batch will get the remaining samples.

- Assume that data and labels are of the same size.

`def add_elementwise(list1, list2)`: Returns `list3` which is an elementwise of `list1`, `list2`.

- Assume `list1`, `list2` is of the same size.

Note – each function of the above can be implemented in one line.

Now check the following:


Run `main.py`, and make sure the neural network is training as supposed.

Make sure the final accuracy is above 95%.

Note: You can adjust the learning rate in `main.py`.



Part 2



In this part you will see how we can achieve a significant speed-up using the GPU.

You will implement a function that calculates the element-wise maximum between two larger scale matrixes.

Given matrixes A, B of size $(1000, 1000)$ with integer values of range $[0, 255]$, the function should return a matrix C of the same size, where $C_{ij} = \max\{A_{ij}, B_{ij}\}$.

Implement the following functions in the file `max_functions.py`:

`def max_cpu (A, B):` Calculates the element-wise maximum on the CPU and returns it.

- Do not use numpy vectorize operations.

`def max_numba (A, B):` Use the NJIT to speed up the above calculation.

`def max_gpu (A, B):` Calculates the element-wise maximum between A, B on the GPU, by invoking the `max_kernel` function with 1000 block, where each block contains 1000 threads.

- `max_kernel` is defined in the same file.

Now do the following:

Run `max_functions.py` on 1 core (flag `-c 1`) to see time comparisons.

Make sure that the NUMBA and GPU calculations are correct.

Include a screenshot of the time comparison between the three methods, and an explanation about the GPU implementation, in the report to be followed.



In addition, run the `max_functions.py` with 2, 4, and 8 cores, and explain the difference.


Notes:

1. You must get a speed-up of at least 40, between the CPU run-time to the GPU run-time.
2. The `max_cpu` function should be implemented in a trivial way. Do not insert trivial delays.
3. You must implement the functions yourself – don't include any existing functions from `numpy` or any other library.
4. You may use `cuda atomic add`, `cuda atomic max` and `cuda syncthreads`.



Part 3

In this part you will implement a matrix calculation between two matrixes.



Specifically, we are interested in a function that given matrix X will calculate $X \cdot X^T$ efficiently.

Implement the following functions, in `matmul_functions.py`:

`def matmul_transpose_trivial(X)`: Calculates $X \cdot X^T$ in the most trivial way – using 3 nested for loops.

`def matmul_transpose_numba(X)`: Use NJIT to speed up the function from above.

`def matmul_transpose_gpu(X)`: Calculates $X \cdot X^T$ on the GPU.

- You should implement `matmul_kernel` and use it.
- `matmul_kernel` should always be called with 1 thread block which contains 1024 threads.

Run `matmul_functions.py`, which will generate comparisons of the run-time of the functions above.



Notes

Report

You must include a report.

1. Provide a detailed explanation of your `max_kernel` implementation, include screenshot, and calculate the speedup between `max_gpu/max_numba` and `max_gpu/max_cpu`, and explanation of the results.
2. Provide a detailed explanation of your `matmul_kernel` implementation, include a screenshot and explanation of the results.

Notes and Tips

- You can add variables and prints as you need, but your code must be clear and organized.
- Don't remove prints or comments already in the code, adhere to instruction comments.
- Document your code thoroughly.

Server

Full explanation can be found in the Jupyter notebook at the course website (in HW1 section).

Submission

Submit a `hw1.zip` with the following files only:

- `utils.py` with your implementation.
- `max_functions.py` with your implementations.
- `matmul_functions.py` with your implementations.
- `hw1.pdf` report of performance analysis.